The Relational Data Model
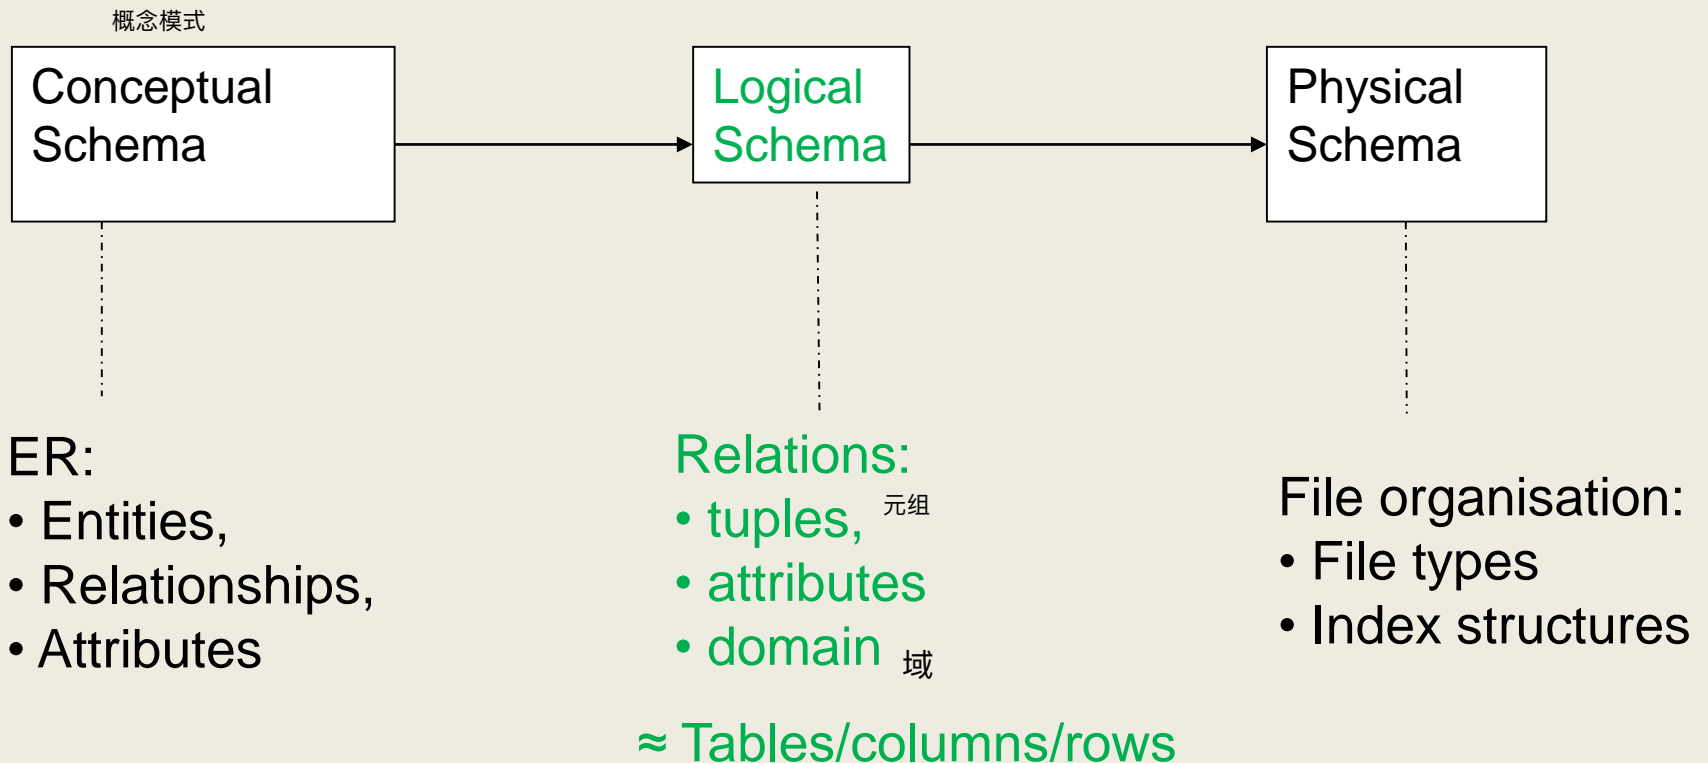
Textbook: chapter 5

Relational Data Model

Different schemas are based on different levels of abstraction

概念模式

| Conceptual Schema | → | Logical Schema | → | Physical Schema |

ER:
- Entities,
- Relationships,
- Attributes

Relations:
- tuples, 元组
- attributes
- domain 域

≈ Tables/columns/rows

File organisation:
- File types
- Index structures

Relational Data Model Concepts

The relational data model is the most widely used data model for database systems.

The *relational data model* describes the world as
- a **collection** of inter-connected ***relations***

Goal of relational model:
- a simple, general data modelling formalism
- which maps easily to file structures (i.e. implementable)

Relational model has **two styles** of terminology:
- mathematical: relation, tuple, attribute, ...
- data-oriented: table, record, field/column, ...

Warning: textbooks alternate between the two; treat them as synonyms

Structures

In the relational model, everything is described using relations.

A relation can be thought of as a named table.

--- Each column of the table corresponds to a named attribute.
列

--- Each row of the table is called a tuple of the relation.
行

The set of allowed values for an attribute is called its domain.

N.B. There is no ordering of column or rows.

Example

PLAYER					
Name	Position	Goals	Age	Height	Weight
Heady	Half-forward	17	24	183	83
Sumich	Full-forward	59	26	191	92
Langdon	Utility	23	23	189	86

PLAYER					
Name	Age	Height	Weight	Goals	Position
Sumich	26	191	92	59	Full-forward
Langdon	23	189	86	23	Utility
Heady	24	183	83	17	Half-forward

Above two tables are the same relation ---- Player

Relational Data Model

Mathematically,

单一元素

- ◦ a *domain D* is a set of atomic values (having some fixed data type) which represent some semantic meaning.

- ◦ an *attribute A* is the name of a role played by a *domain*, *dom(A)*.

- ◦ a *relation schema R*, denoted by

$R(A_1, A_2, ..., A_n)$, is a set of attributes

$R = \{A_1, A_2, ..., A_n\}$.

Composite and multivalued attributes are disallowed!

Why Relational Model?

- Very simple model

- Often a good match for the way we think about our data

- Foundations in logic and set theory (will introduce in later parts of the course)

- Abstract model that underlies SQL, the most important language in DBMSs today

Relational Data Model vs ER Model

Relation schema (intension) ⇄ entity or relationship type schema (intension).

attributes ⇄ attributes

tuple ⇄ instance of entity/relationship

relation (instance, extension) ⇄ entity/relationship extension

composite and multivalued attributes are allowed in ER model, but not allowed in relational data model.

Keys

Keys are used to identify tuples in a relation.

A *superkey* is a set of attributes that uniquely determines a tuple.

Note that this is a property of the relation that does not depend on the current relation instance.
关系实例

A *candidate key* is a superkey, none of whose *proper* subsets is a superkey.
候选码/键   真子集

Keys are determined by the applications.

Keys

PLAYER					
Name	Position	Goals	Age	Height	Weight
Heady	Half-forward	17	24	183	83
Sumich	Full-forward	59	26	191	92
Langdon	Utility	23	23	189	86

E.g. if {Name} is unique then it is a candidate key for PLAYER; otherwise we need to use the whole tuple or create a candidate key, say PID.

{Goals} usually cannot not be a candidate key since different players *might* have the same number of goals.

{Name, Goals} is a superkey but not a candidate key (because {Name} is a key).

(*Assumption here is no two players have the same name*)

Keys

PLAYER						
Person_ID	Name	Position	Goals	Age	Height	Weight
1	Heady	Half-forward	17	24	183	83
2	Sumich	Full-forward	59	26	191	92
3	Langdon	Utility	23	23	189	86

A *primary key* is a designated candidate key.
主键   指定的

In many applications it is necessary to invent a primary key if there is no natural one - often this would be a non-negative integer

e.g. Person_ID.

When a relation schema has several candidate keys, usually better to choose a primary key with a single attribute or a small number of attributes.

Integrity constraints

There are several kinds of integrity constraints that are an integral part of the relational model:

Key constraint: candidate key values must be unique for every relation instance.

Entity integrity: an attribute that is part of a primary key cannot be NULL.

Referential integrity: The third kind has to do with "foreign keys".

参照完整性
引用完整性

Keys

Foreign keys are used to refer to a tuple in another relation.

A set, *FK*, of attributes from a relation schema R1 may be a foreign *key* if

- the attributes have the same domains as the attributes in the primary key of another

  域

  relation schema R_2, and

- a value of FK in a tuple t_1 of R_1 either occurs as a value of PK for some tuple t_2 in R_2
  or is null.

Referential integrity: The value of *FK* must occur in the other relation or be
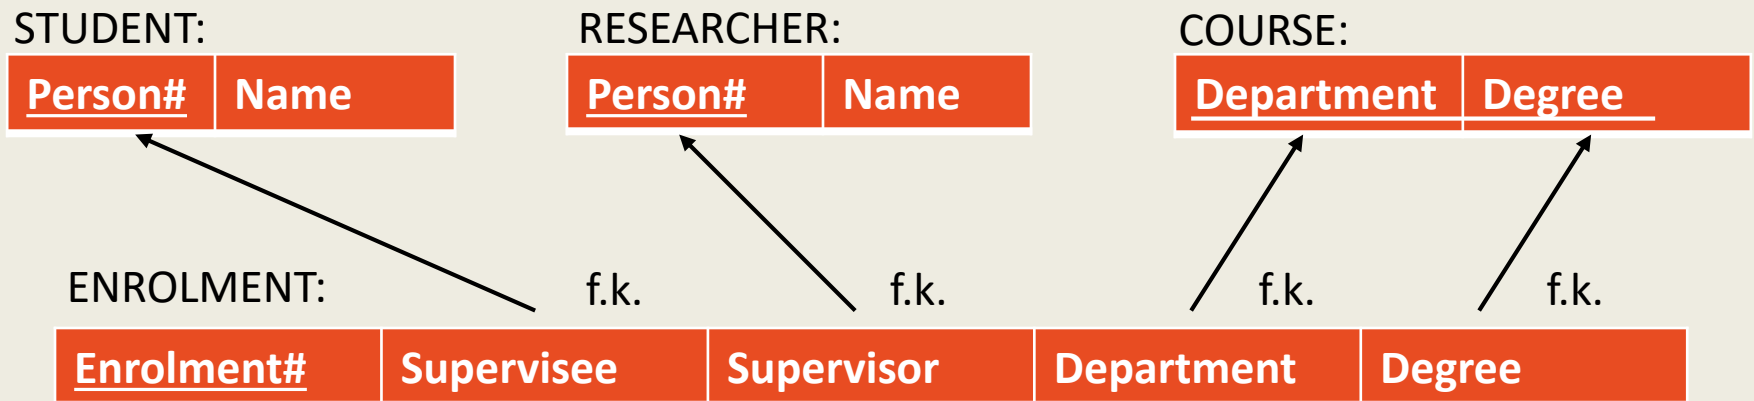
entirely NULL.

Referential Integrity

Foreign keys are critical in relational DBs; they provide …

- the "glue" that links individual relations (we will see more about this in SQL part of the course)
- the way to assemble query answers from multiple tables
- the relational representation of ER relationships

Checking constraints on updates

– To maintain the integrity of the database, we need to check that integrity constraints will not be violated before proceeding with an update.

– Example: Suppose we have the following schema with foreign keys as shown:

STUDENT:

Person#	Name

RESEARCHER:

Person#	Name

COURSE:

Department	Degree

ENROLMENT: f.k.   f.k.   f.k.   f.k.

Enrolment#	Supervisee	Supervisor	Department	Degree

STUDENT:

Person#	Name
1	Dr C.C.Chen
3	Ms K.Juliff
4	Ms J.Gledill
5	Ms B.K.Lee

RESEARCHER:

Person#	Name
1	Dr C.C.Chen
2	Dr R.G.Wilkinson

COURSE:

Department	Degree
Psychology	Ph.D.
Comp.Sci.	Ph.D.
Comp.Sci.	M.Sc.
Psychology	M.Sc.

ENROLMENT:

Enrolment#	Supervisee	Supervisor	Department	Degree
1	1	2	Psychology	Ph.D.
2	3	1	Comp.Sci.	Ph.D.
3	4	1	Comp.Sci.	M.Sc.
4	5	1	Comp.Sci.	M.Sc.

Insertions

Insertions: When inserting, we need to check

- that the candidate keys are not already present,

- that the value of each foreign key either

    – is all null, or

    – is all non-NULL and occurs in the referenced relation.

STUDENT:

Person#	Name
1	Dr C.C.Chen
3	Ms K.Juliff
4	Ms J.Gledill
5	Ms B.K.Lee

RESEARCHER:

Person#	Name
1	Dr C.C.Chen
2	Dr R.G.Wilkinson

COURSE:

Department	Degree
Psychology	Ph.D.
Comp.Sci.	Ph.D.
Comp.Sci.	M.Sc.
Psychology	M.Sc.

ENROLMENT:

Enrolment#	Supervisee	Supervisor	Department	Degree
1	1	2	Psychology	Ph.D.
2	3	1	Comp.Sci.	Ph.D.
3	4	1	Comp.Sci.	M.Sc.
4	5	1	Comp.Sci.	M.Sc.

1. Insert $< 2, Dr.V.Ciesielski >$ into RESEARCHER

STUDENT:

Person#	Name
1	Dr C.C.Chen
3	Ms K.Juliff
4	Ms J.Gledill
5	Ms B.K.Lee

RESEARCHER:

Person#	Name
1	Dr C.C.Chen
2	Dr R.G.Wilkinson

COURSE:

Department	Degree
Psychology	Ph.D.
Comp.Sci.	Ph.D.
Comp.Sci.	M.Sc.
Psychology	M.Sc.

ENROLMENT:

Enrolment#	Supervisee	Supervisor	Department	Degree
1	1	2	Psychology	Ph.D.
2	3	1	Comp.Sci.	Ph.D.
3	4	1	Comp.Sci.	M.Sc.
4	5	1	Comp.Sci.	M.Sc.

1. Insert $< 2, Dr.V.Ciesielski >$ into RESEARCHER
Allowed? No. Violates a key constraint.
Action? Reject or allow the user to correct.

STUDENT:

Person#	Name
1	Dr C.C.Chen
3	Ms K.Juliff
4	Ms J.Gledill
5	Ms B.K.Lee

RESEARCHER:

Person#	Name
1	Dr C.C.Chen
2	Dr R.G.Wilkinson

COURSE:

Department	Degree
Psychology	Ph.D.
Comp.Sci.	Ph.D.
Comp.Sci.	M.Sc.
Psychology	M.Sc.

ENROLMENT:

Enrolment#	Supervisee	Supervisor	Department	Degree
1	1	2	Psychology	Ph.D.
2	3	1	Comp.Sci.	Ph.D.
3	4	1	Comp.Sci.	M.Sc.
4	5	1	Comp.Sci.	M.Sc.

2. Insert $< Comp.Sci., NULL >$ into COURSE
Allowed?

STUDENT:

Person#	Name
1	Dr C.C.Chen
3	Ms K.Juliff
4	Ms J.Gledill
5	Ms B.K.Lee

RESEARCHER:

Person#	Name
1	Dr C.C.Chen
2	Dr R.G.Wilkinson

COURSE:

Department	Degree
Psychology	Ph.D.
Comp.Sci.	Ph.D.
Comp.Sci.	M.Sc.
Psychology	M.Sc.

ENROLMENT:

Enrolment#	Supervisee	Supervisor	Department	Degree
1	1	2	Psychology	Ph.D.
2	3	1	Comp.Sci.	Ph.D.
3	4	1	Comp.Sci.	M.Sc.
4	5	1	Comp.Sci.	M.Sc.

2. Insert < *Comp.Sci.,NULL* > into COURSE
   Allowed? No. Violates the entity integrity constraint.
   Action: Reject or correct.

STUDENT:

Person#	Name
1	Dr C.C.Chen
3	Ms K.Juliff
4	Ms J.Gledill
5	Ms B.K.Lee

RESEARCHER:

Person#	Name
1	Dr C.C.Chen
2	Dr R.G.Wilkinson

COURSE:

Department	Degree
Psychology	Ph.D.
Comp.Sci.	Ph.D.
Comp.Sci.	M.Sc.
Psychology	M.Sc.

ENROLMENT:

Enrolment#	Supervisee	Supervisor	Department	Degree
1	1	2	Psychology	Ph.D.
2	3	1	Comp.Sci.	Ph.D.
3	4	1	Comp.Sci.	M.Sc.
4	5	1	Comp.Sci.	M.Sc.

3. Insert < 5, 6, 2, *Psychology, Ph.D.* > into ENROLMENT
   Allowed?

STUDENT:

Person#	Name
1	Dr C.C.Chen
3	Ms K.Juliff
4	Ms J.Gledill
5	Ms B.K.Lee

RESEARCHER:

Person#	Name
1	Dr C.C.Chen
2	Dr R.G.Wilkinson

COURSE:

Department	Degree
Psychology	Ph.D.
Comp.Sci.	Ph.D.
Comp.Sci.	M.Sc.
Psychology	M.Sc.

ENROLMENT:

Enrolment#	Supervisee	Supervisor	Department	Degree
1	1	2	Psychology	Ph.D.
2	3	1	Comp.Sci.	Ph.D.
3	4	1	Comp.Sci.	M.Sc.
4	5	1	Comp.Sci.	M.Sc.

3. Insert < 5, 6, 2, *Psychology, Ph.D.* > into ENROLMENT
   Allowed? No. Violates a referential integrity constraint (There is no person number 6).
   Action: Reject, correct or accept after insertion of person number 6.

Deletions

Deletions: When deleting, we need to check referential integrity – check whether the primary key occurs in another relation.

Examples:

RESEARCHER:

Person#	Name
1	Dr C.C.Chen
2	Dr R.G.Wilkinson

1. Delete tuple with Person# = 2 from RESEARCHER

   Allowed? No. Violates the referential integrity.

   Action: Reject, correct or modify the ENROLMENT tuple by the following actions:

Deletions

deleting it (note that the this requires another integrity check, possibly causing a cascade of deletions), or

setting the foreign key value to NULL (note this can't be done if it is part of a primary key), or

setting the foreign key value to another

  acceptable value.

Modifications 修改

If the modified attribute is a

- primary key: this is similar to deleting and then reinserting.

- foreign key: check that the new value refers to an existing tuple.

- neither: no problems can arise.

Relational database definition

◦ A *relational database schema*, is a set of relation schema $\{R_1, \ldots, R_m\}$ and a set of integrity constraints.完整性约束

◦ A relational database instance is a set of relation instances $\{r_1, \ldots, r_m\}$ such that each r_i is an instance of R_i, and the integrity constraints are satisfied.

ER to Relational Data Model Mapping

One technique for database design is to first design a conceptual schema using a high-level data model, and then map it to a conceptual schema in the DBMS data model for the chosen DBMS.

Here we look at a way to do this mapping from the ER to the relational data model.

It involves the following 7 steps.

Example

SSN

Fname

Lname

Birdate

Works_For

Phone

Name

Location

m

Mdate

Employee

Department

1

1

1

Manager

1

Family

m

Participation

n

Project

m

Dependent

Time

Fname

Birdate

Relation

Pname

Pnumber

29

ER to RDM

Step 1 : For each regular (not weak) entity type E, create a relation R with

- Attributes : All simple attributes (and simple components of composite attributes) of E.

- Key : Choose one of the keys of E as the primary key for the relation. In our example, **Pname** is a candidate key for **Project** relation, **Pnumber** is also a candidate key for **Project** relation, we choose **Pname** as the primary key.

ER to RDM

Step 1a : For each specialised entity type E, with parent entity type P, create a relation R with

- Attributes : The attributes of the key of P, plus the simple attributes of E.

- Key : The key of P.
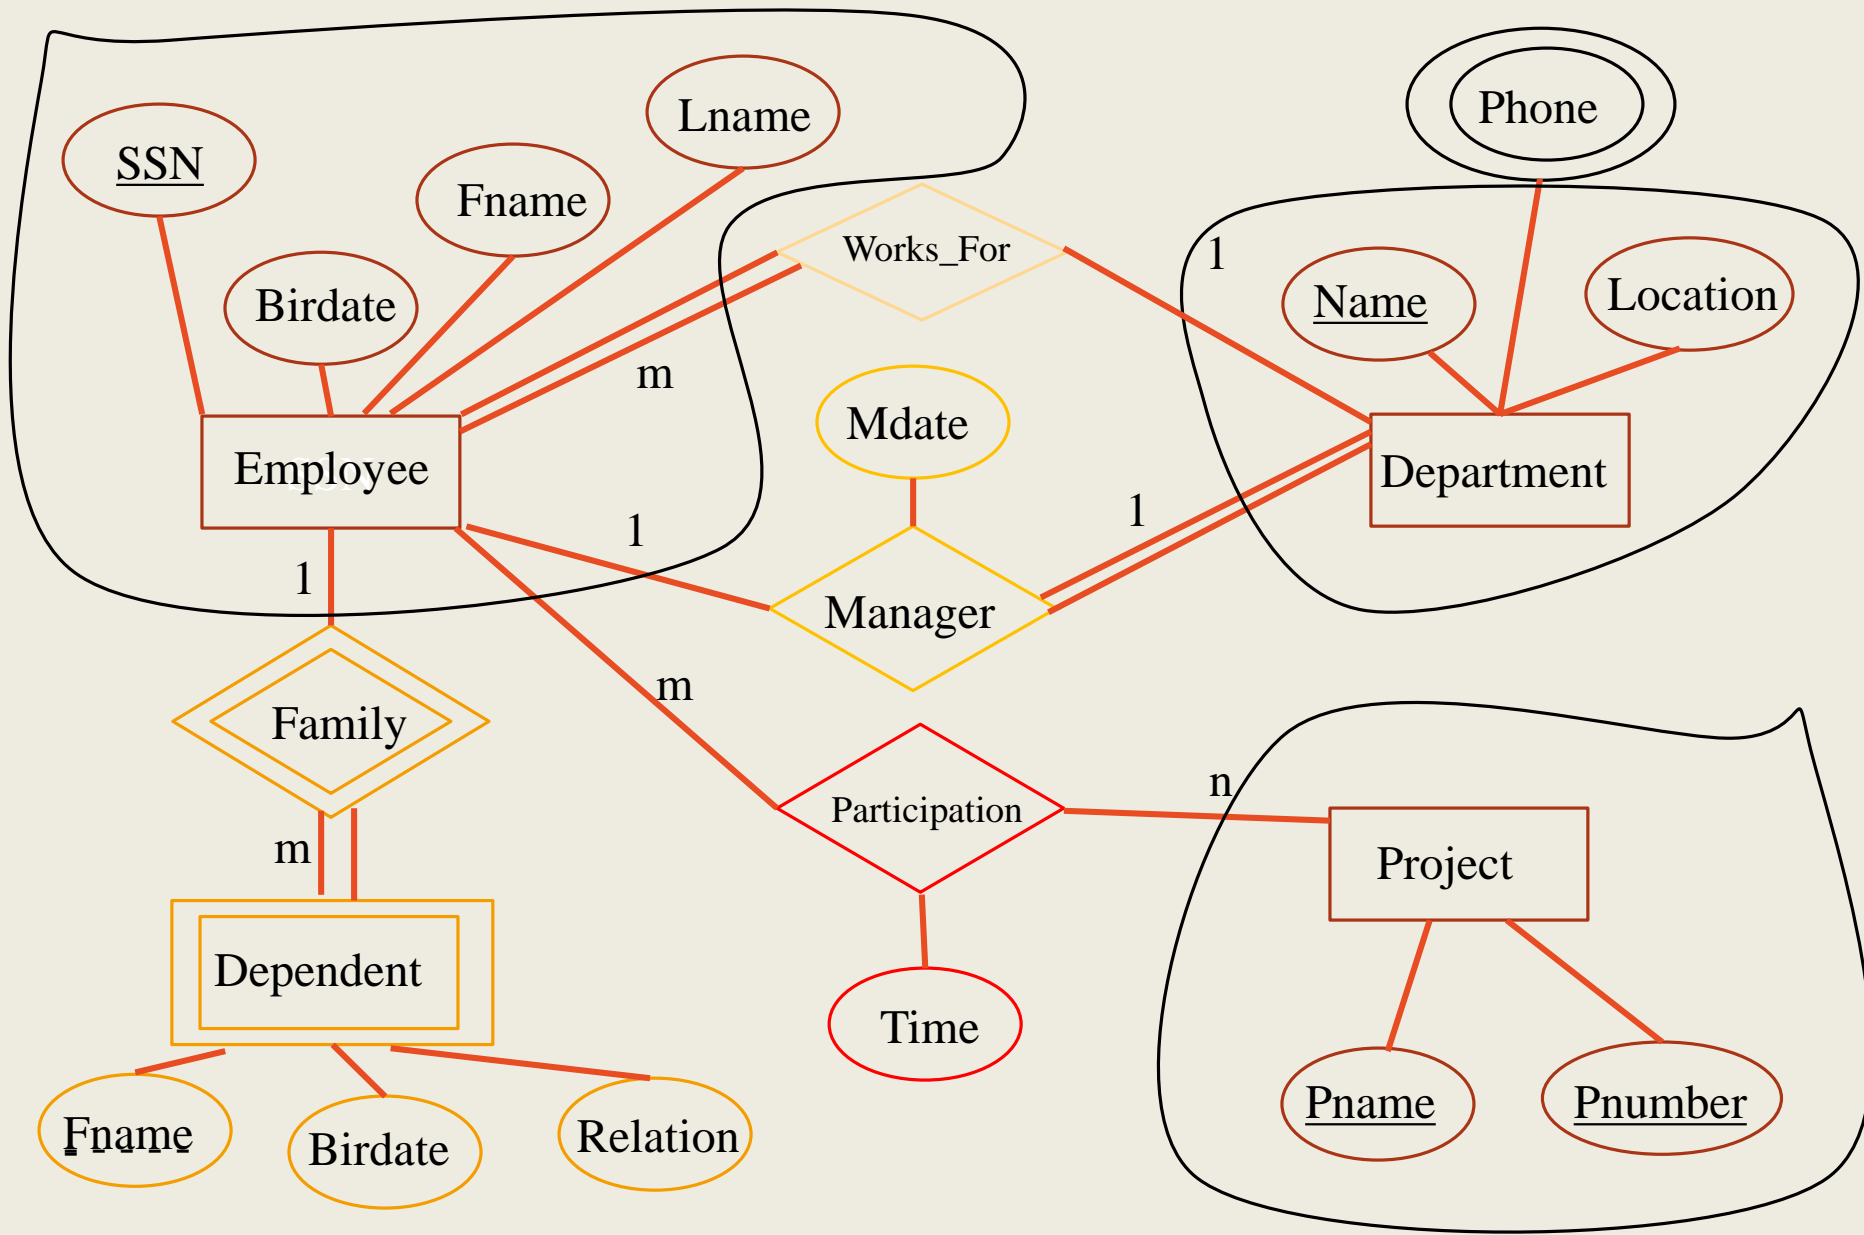
Employee

SSN	Fname	Lname	Birdate

Department

Name	Location

Project

Pname	Pnumber

ER to RDM

Step 2 : For each weak entity type W, with owner entity type E, create a relation R with

- Attributes : All simple attributes (and simple components of composite attributes) of W, and include as a foreign key the prime attributes of the relation derived from E.

- Key : The foreign key plus the partial key of W.

Employee

SSN	Fname	Lname	Birdate

Department

Name	Location

Project

Pname	Pnumber

Dependent

SSN	Fname	Birdate	Relation

ER to RDM

Step 3 : For each 1:1 relationship type B. Let E and F be the participating entity types. Let S and T be the corresponding relations.

- Choose one of S and T (prefer one that participates totally), say S.

- Add the attributes of the primary key of T to S as a foreign key.

- Add the simple attributes (and simple components of composite attributes) of B as attributes of S.

(Alternative: merge the two entity types and the relationship into a single relation, especially if both participate totally and do not participate in other relationships).

ER to RDM

Step 4 : For each regular 1:N relationship type B.

– Let E and F be the participating entity types.

– Let E by the entity type on the 1 side, F the one on the N side.

– Let S and T be the corresponding relations.

– Add the attributes of the primary key of S to T as a foreign key.

– Add to T any simple attributes (or simple components of composite attributes) of the relationship.

(Notice that this doesn't add any new tuples, just attributes.)

ER Diagram with the following entities and relationships:

- **Employee** (entity) with attributes: SSN (key), Birdate, Fname, Lname
- **Department** (entity) with attributes: Name (key), Phone (multivalued), Location
- **Works_For** relationship between Employee (m) and Department (1)
- **Manager** relationship between Employee (1) and Department (1), with attribute Mdate
- **Family** (identifying relationship) between Employee (1) and Dependent (m)
- **Dependent** (weak entity) with attributes: Fname, Birdate, Relation
- **Participation** relationship between Employee (m) and Project (n), with attribute Time
- **Project** (entity) with attributes: Pname (key), Pnumber (key)

41

Employee

SSN	Fname	Lname	Birdate	Dname

Department

Name	Location	MSSN	Mdate

Project

| Pname	Pnumber

Dependent

SSN	Fname	Birdate	Relation

ER to RDM

Step 5 : For each N:M relationship type B. Create a new relation R. Let E and F be the participating entity types. Let S and T be the corresponding relations.

- Attributes : The key of S and the key of T as foreign keys, plus the simple attributes (and simple components of composite attributes) of B.
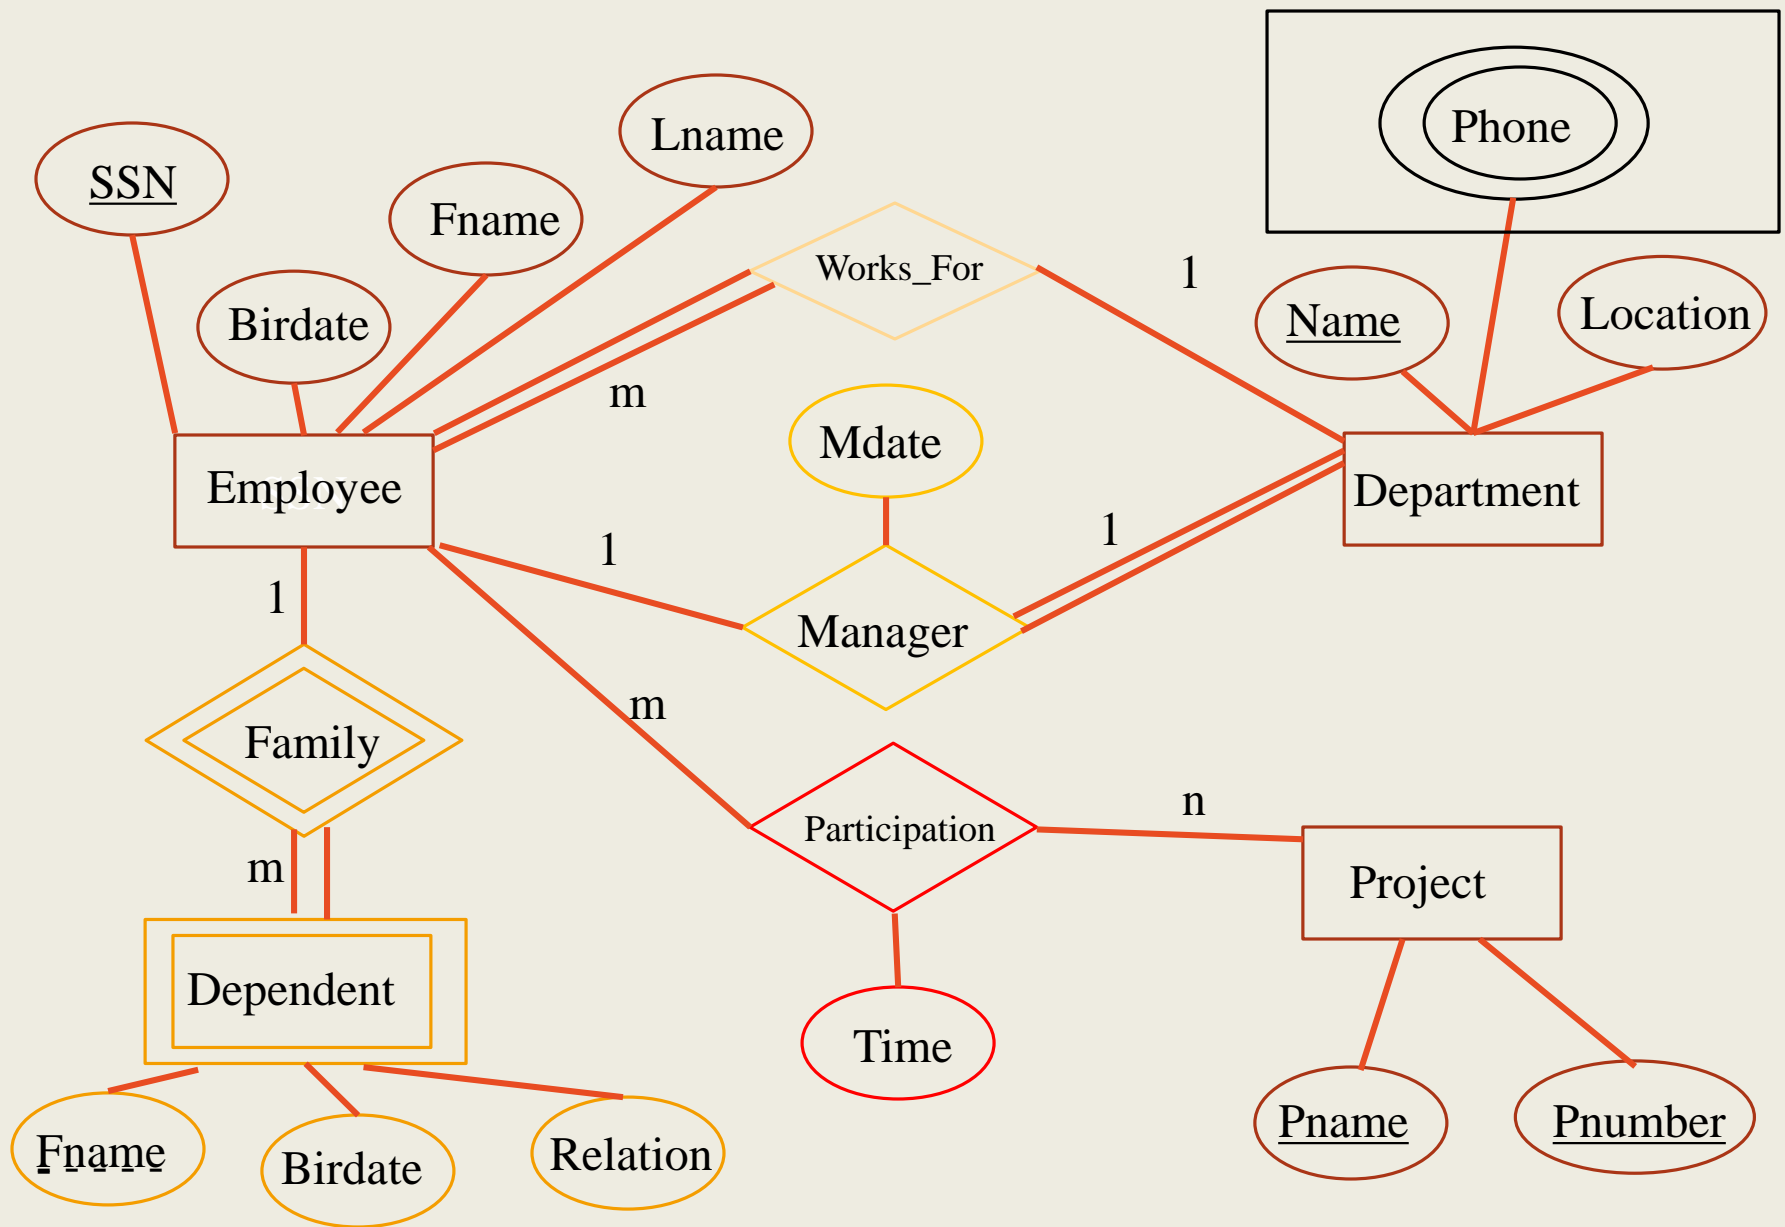
- Key : The key of S and the key of T.
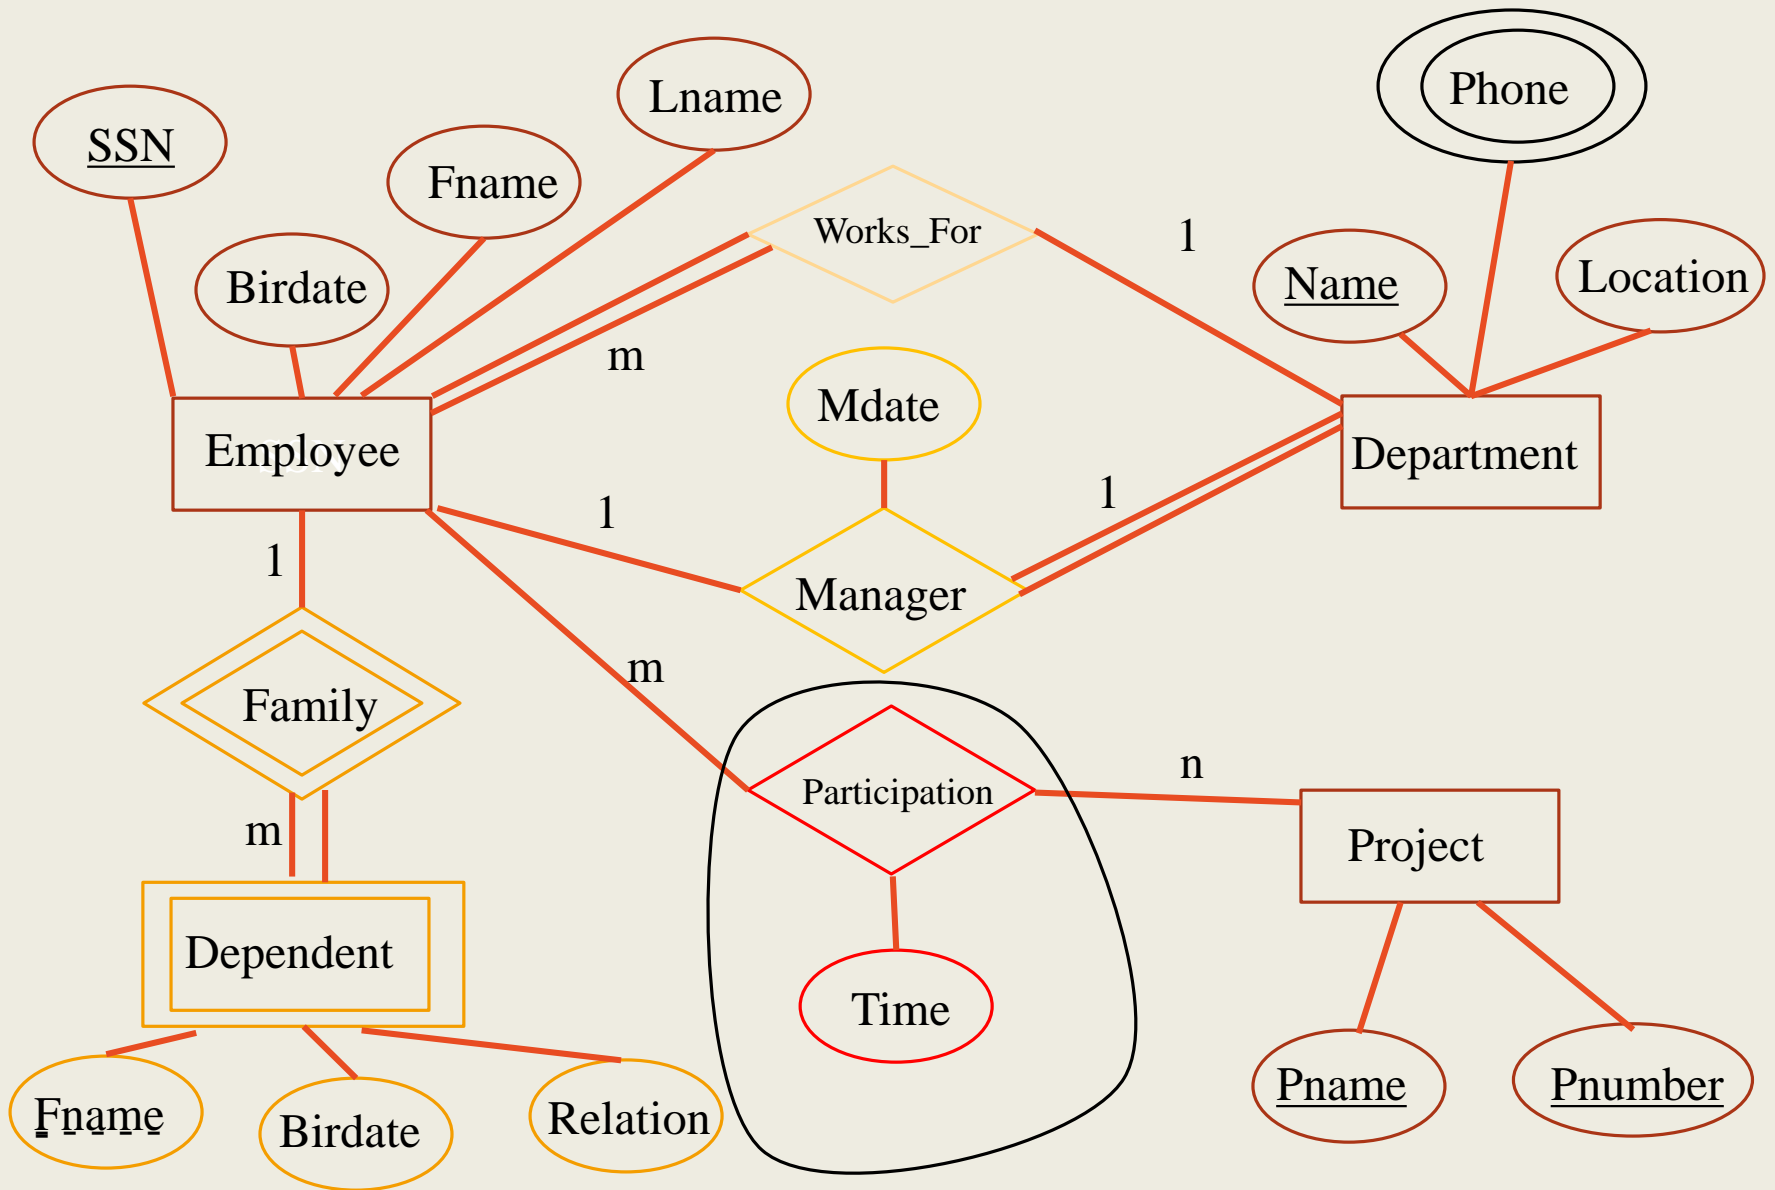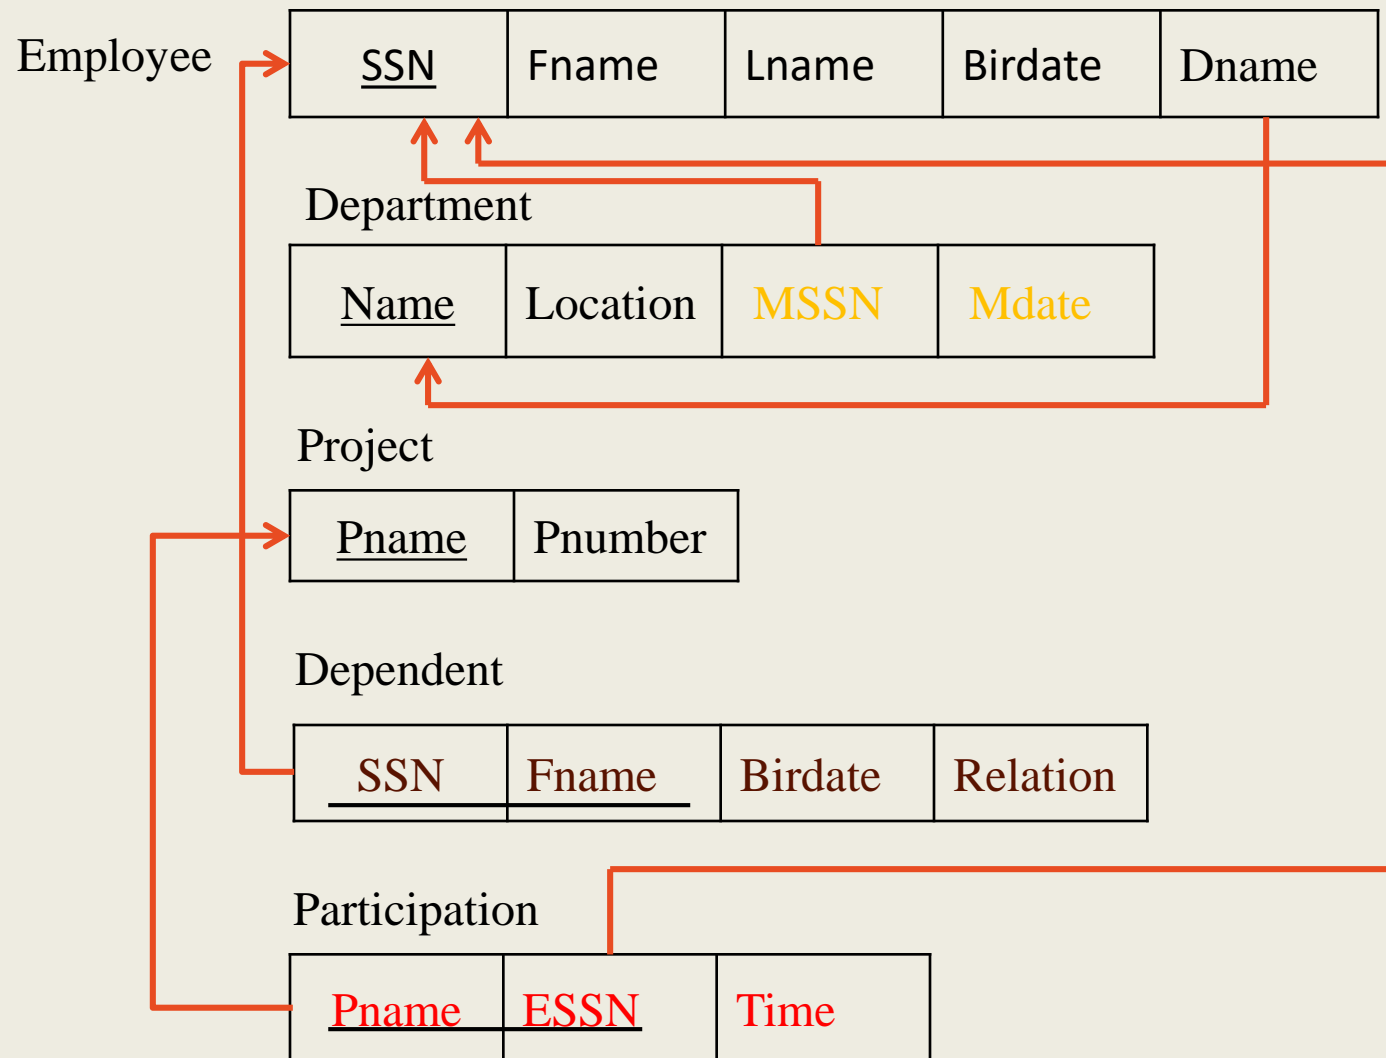
ER to RDM

Step 6 : For each multivalued attribute A. Create a new relation R. Let A be an attribute of E.

– Attributes :

    1. A (if A is a simple attribute) together with the key of

       E as a foreign key.

    2. The simple components of A (if A is a composite

       attribute), together with the key of E as a foreign key.

– Key : All attributes.

Employee

SSN	Fname	Lname	Birdate	Dname

Department

Name	Location	MSSN	Msdate

Project

Pname	Pnumber

Dependent

SSN	Fname	Birdate	Relation

Participation

Pname	ESSN	Time

D_number

Dname	Dnumber

48
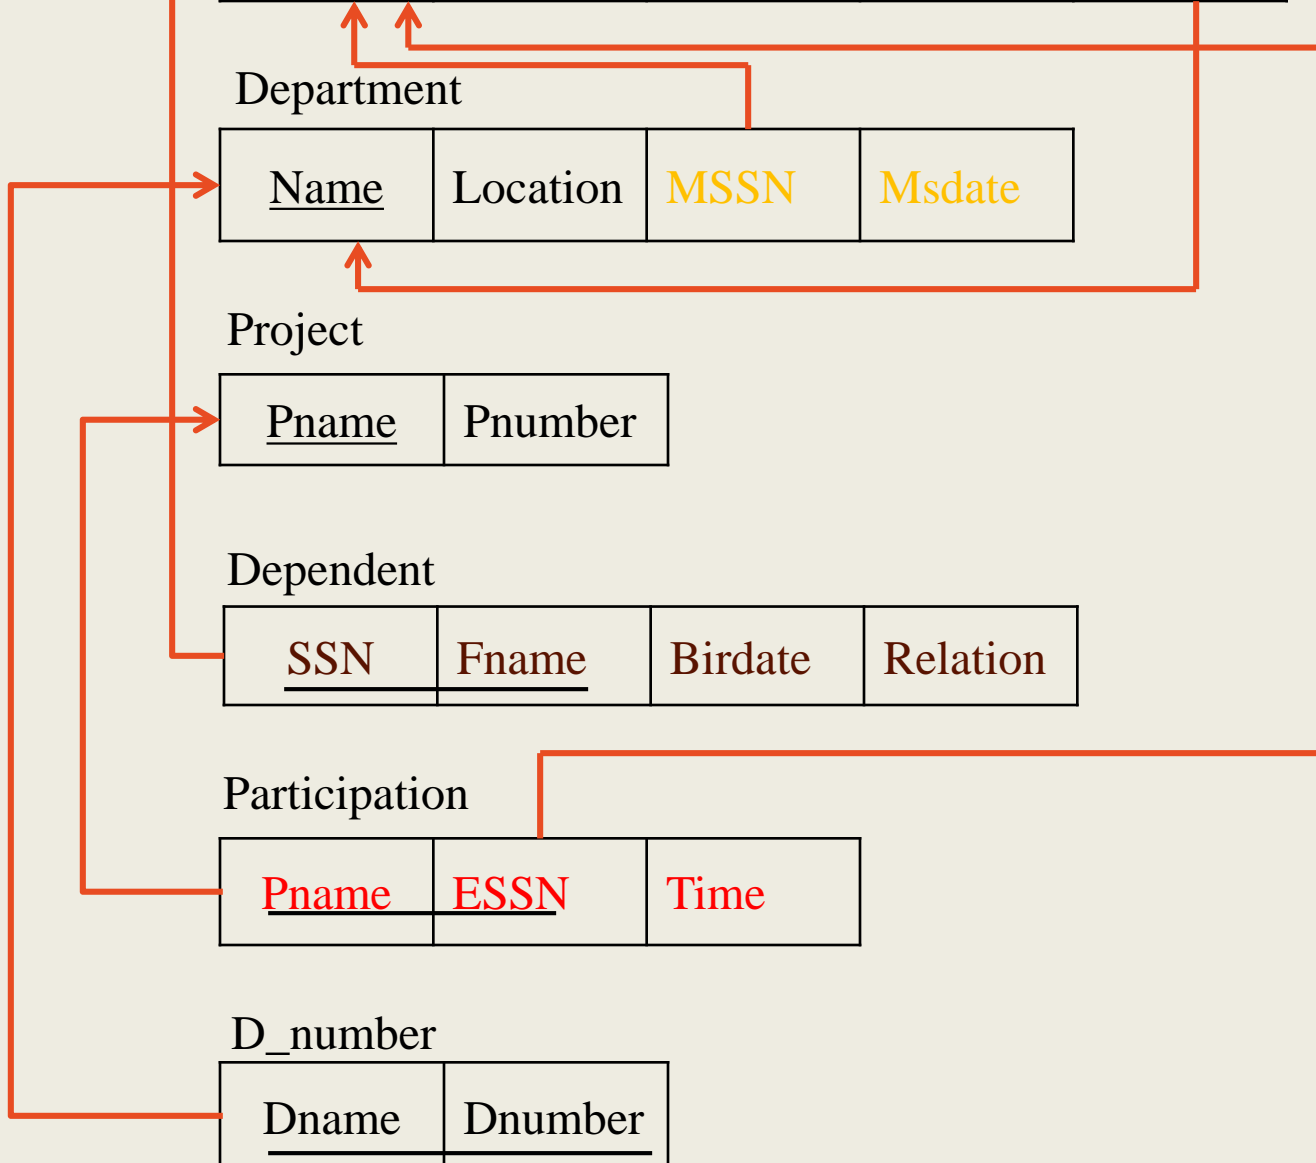
ER to RDM

Step 7 : For each n-ary relationship type (n > 2). Create a new relation with

- Attributes : as for Step 5.

- Key : as for Step 5, except that if one of the participating entity types has participation ratio 1, its key can be used as a key for the new relation.

Learning Outcome

1. Understanding of relational model

2. Given an ERD, convert the ERD to relational model.